

Reshaping and pivot tables

pandas.pydata.org/pandas-docs/stable/user_guide/reshaping.html

Reshaping by pivoting DataFrame objects

Pivot

df

	foo	bar	baz	zoo
0	one	A	1	x
1	one	B	2	y
2	one	C	3	z
3	two	A	4	q
4	two	B	5	w
5	two	C	6	t

```
df.pivot(index='foo',  
          columns='bar',  
          values='baz')
```

bar	A	B	C
foo			
one	1	2	3
two	4	5	6

Data is often stored in so-called “stacked” or “record” format:

In [1]: df

Out[1]:

	date	variable	value
0	2000-01-03	A	0.469112
1	2000-01-04	A	-0.282863
2	2000-01-05	A	-1.509059
3	2000-01-03	B	-1.135632
4	2000-01-04	B	1.212112
5	2000-01-05	B	-0.173215
6	2000-01-03	C	0.119209
7	2000-01-04	C	-1.044236
8	2000-01-05	C	-0.861849
9	2000-01-03	D	-2.104569
10	2000-01-04	D	-0.494929
11	2000-01-05	D	1.071804

For the curious here is how the above `DataFrame` was created:

```
import pandas.util.testing as tm
```

```
tm.N = 3
```

```
def unpivot(frame):
    N, K = frame.shape
    data = {'value': frame.to_numpy().ravel('F'),
            'variable': np.asarray(frame.columns).repeat(N),
            'date': np.tile(np.asarray(frame.index), K)}
    return pd.DataFrame(data, columns=['date', 'variable', 'value'])
```

```
df = unpivot(tm.makeTimeDataFrame())
```

To select out everything for variable `A` we could do:

```
In [2]: df[df['variable'] == 'A']
```

```
Out[2]:
```

	date	variable	value
0	2000-01-03	A	0.469112
1	2000-01-04	A	-0.282863
2	2000-01-05	A	-1.509059

But suppose we wish to do time series operations with the variables. A better representation would be where the `columns` are the unique variables and an `index` of dates identifies individual observations. To reshape the data into this form, we use the `DataFrame.pivot()` method (also implemented as a top level function `pivot()`):

```
In [3]: df.pivot(index='date', columns='variable', values='value')
```

```
Out[3]:
```

variable	A	B	C	D
date				
2000-01-03	0.469112	-1.135632	0.119209	-2.104569
2000-01-04	-0.282863	1.212112	-1.044236	-0.494929
2000-01-05	-1.509059	-0.173215	-0.861849	1.071804

If the `values` argument is omitted, and the input `DataFrame` has more than one column of values which are not used as column or index inputs to `pivot`, then the resulting “pivoted” `DataFrame` will have `hierarchical columns` whose topmost level indicates the respective value column:

```
In [4]: df['value2'] = df['value'] * 2
```

```
In [5]: pivoted = df.pivot(index='date', columns='variable')
```

```
In [6]: pivoted
```

```
Out[6]:
```

	value				value2			
variable	A	B	C	D	A	B	C	D
date								
2000-01-03	0.469112	-1.135632	0.119209	-2.104569	0.938225	-2.271265	0.238417	-4.209138
2000-01-04	-0.282863	1.212112	-1.044236	-0.494929	-0.565727	2.424224	-2.088472	-0.989859
2000-01-05	-1.509059	-0.173215	-0.861849	1.071804	-3.018117	-0.346429	-1.723698	2.143608

You can then select subsets from the pivoted `DataFrame` :

```
In [7]: pivoted['value2']
```

```
Out[7]:
```

variable	A	B	C	D
date				
2000-01-03	0.938225	-2.271265	0.238417	-4.209138
2000-01-04	-0.565727	2.424224	-2.088472	-0.989859
2000-01-05	-3.018117	-0.346429	-1.723698	2.143608

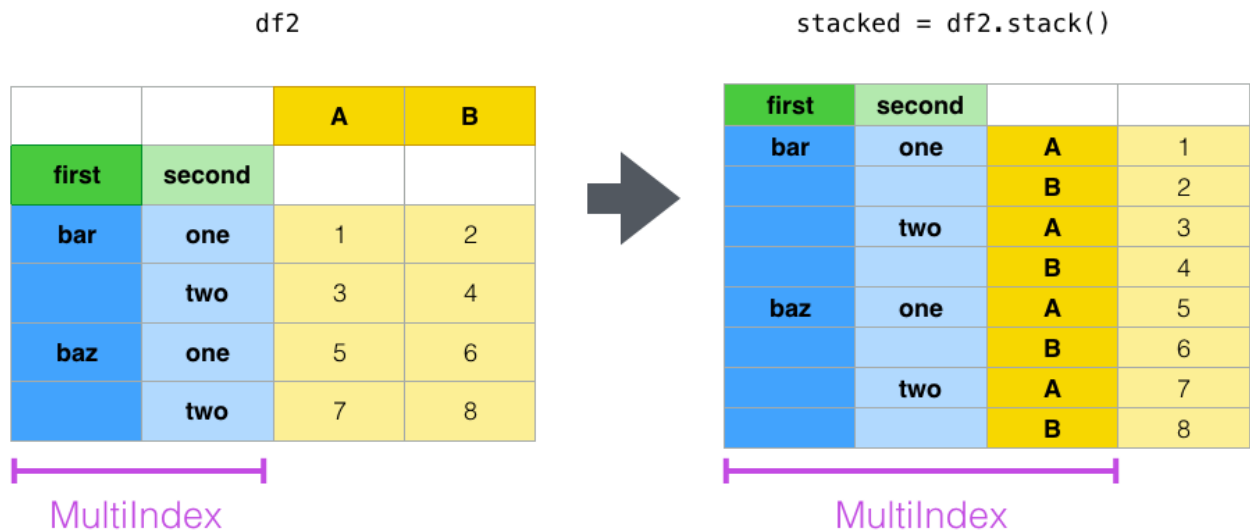
Note that this returns a view on the underlying data in the case where the data are homogeneously-typed.

Note

`pivot()` will error with a `ValueError: Index contains duplicate entries, cannot reshape` if the index/column pair is not unique. In this case, consider using `pivot_table()` which is a generalization of pivot that can handle duplicate values for one index/column pair.

Reshaping by stacking and unstacking

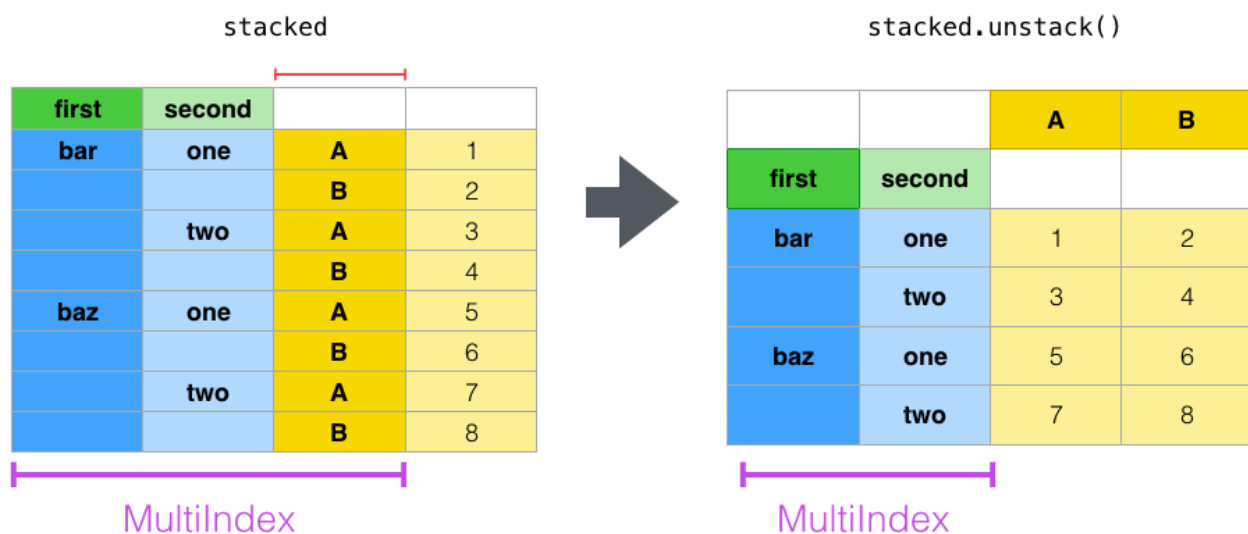
Stack



Closely related to the `pivot()` method are the related `stack()` and `unstack()` methods available on `Series` and `DataFrame`. These methods are designed to work together with `MultiIndex` objects (see the section on [hierarchical indexing](#)). Here are essentially what these methods do:

- `stack`: “pivot” a level of the (possibly hierarchical) column labels, returning a `DataFrame` with an index with a new inner-most level of row labels.
- `unstack`: (inverse operation of `stack`) “pivot” a level of the (possibly hierarchical) row index to the column axis, producing a reshaped `DataFrame` with a new inner-most level of column labels.

Unstack



The clearest way to explain is by example. Let's take a prior example data set from the hierarchical indexing section:

```
In [8]: tuples = list(zip(*[['bar', 'bar', 'baz', 'baz',
...:                        'foo', 'foo', 'qux', 'qux'],
...:                        ['one', 'two', 'one', 'two',
...:                        'one', 'two', 'one', 'two']]))
...:
```

```
In [9]: index = pd.MultiIndex.from_tuples(tuples, names=['first', 'second'])
```

```
In [10]: df = pd.DataFrame(np.random.randn(8, 2), index=index, columns=['A', 'B'])
```

```
In [11]: df2 = df[:4]
```

```
In [12]: df2
```

```
Out[12]:
```

```

      A      B
first second
bar one  0.721555 -0.706771
     two -1.039575  0.271860
baz one  -0.424972  0.567020
     two  0.276232 -1.087401
```

The `stack` function "compresses" a level in the `DataFrame`'s columns to produce either:

- A `Series`, in the case of a simple column Index.
- A `DataFrame`, in the case of a `MultiIndex` in the columns.

If the columns have a `MultiIndex`, you can choose which level to stack. The stacked level becomes the new lowest level in a `MultiIndex` on the columns:

```
In [13]: stacked = df2.stack()
```

```
In [14]: stacked
```

```
Out[14]:
```

```
first second
bar one A 0.721555
      B -0.706771
    two A -1.039575
      B 0.271860
baz one A -0.424972
      B 0.567020
    two A 0.276232
      B -1.087401
dtype: float64
```

With a “stacked” `DataFrame` or `Series` (having a `MultiIndex` as the `index`), the inverse operation of `stack` is `unstack`, which by default unstacks the **last level**:

```
In [15]: stacked.unstack()
```

```
Out[15]:
```

```
      A      B
first second
bar one 0.721555 -0.706771
    two -1.039575 0.271860
baz one -0.424972 0.567020
    two 0.276232 -1.087401
```

```
In [16]: stacked.unstack(1)
```

```
Out[16]:
```

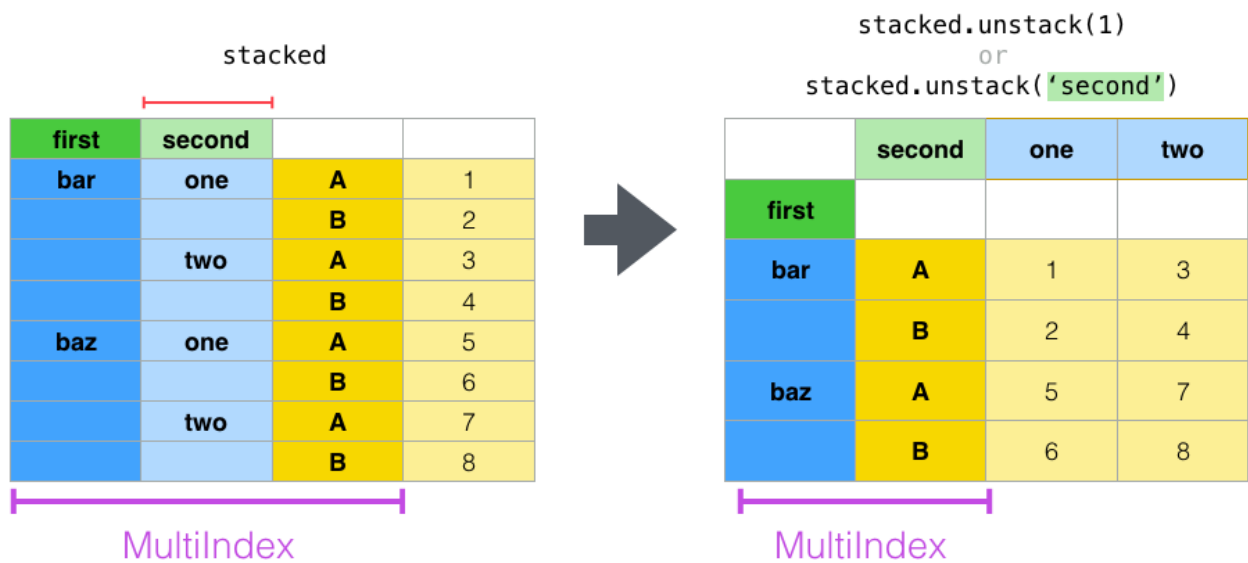
```
second one two
first
bar A 0.721555 -1.039575
    B -0.706771 0.271860
baz A -0.424972 0.276232
    B 0.567020 -1.087401
```

```
In [17]: stacked.unstack(0)
```

```
Out[17]:
```

```
first bar baz
second
one A 0.721555 -0.424972
    B -0.706771 0.567020
two A -1.039575 0.276232
    B 0.271860 -1.087401
```

Unstack(1)



If the indexes have names, you can use the level names instead of specifying the level numbers:

```
In [18]: stacked.unstack('second')
```

```
Out[18]:
```

```
second    one    two
```

```
first
```

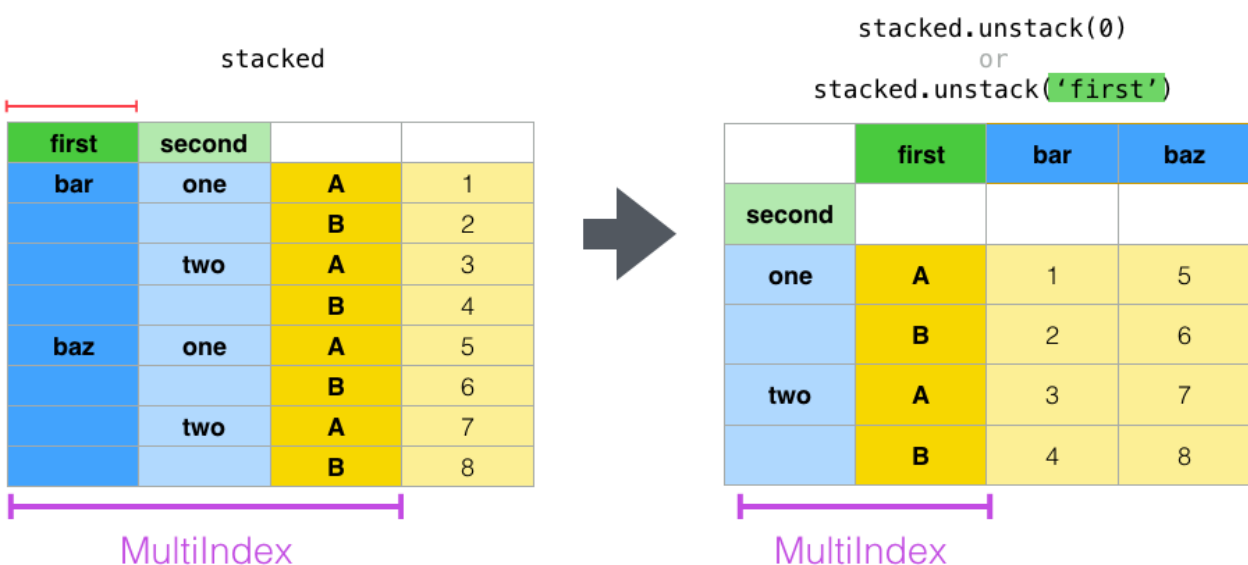
```
bar  A  0.721555 -1.039575
```

```
      B -0.706771  0.271860
```

```
baz  A -0.424972  0.276232
```

```
      B  0.567020 -1.087401
```

Unstack(0)



Notice that the `stack` and `unstack` methods implicitly sort the index levels involved. Hence a call to `stack` and then `unstack`, or vice versa, will result in a **sorted** copy of the original `DataFrame` or `Series` :

```
In [19]: index = pd.MultiIndex.from_product([[2, 1], ['a', 'b']])
```

```
In [20]: df = pd.DataFrame(np.random.randn(4), index=index, columns=['A'])
```

```
In [21]: df
```

```
Out[21]:
```

```
      A
2 a -0.370647
  b -1.157892
1 a -1.344312
  b  0.844885
```

```
In [22]: all(df.unstack().stack() == df.sort_index())
```

```
Out[22]: True
```

The above code will raise a `TypeError` if the call to `sort_index` is removed.

Multiple levels

You may also stack or unstack more than one level at a time by passing a list of levels, in which case the end result is as if each level in the list were processed individually.


```
In [23]: columns = pd.MultiIndex.from_tuples([
.....:     ('A', 'cat', 'long'), ('B', 'cat', 'long'),
.....:     ('A', 'dog', 'short'), ('B', 'dog', 'short')],
.....:     names=['exp', 'animal', 'hair_length']
.....: )
.....:
```

```
In [24]: df = pd.DataFrame(np.random.randn(4, 4), columns=columns)
```

```
In [25]: df
```

```
Out[25]:
```

exp	A	B	A	B
animal	cat	cat	dog	dog
hair_length	long	long	short	short
0	1.075770	-0.109050	1.643563	-1.469388
1	0.357021	-0.674600	-1.776904	-0.968914
2	-1.294524	0.413738	0.276662	-0.472035
3	-0.013960	-0.362543	-0.006154	-0.923061

```
In [26]: df.stack(level=['animal', 'hair_length'])
```

```
Out[26]:
```

exp	A	B
animal hair_length		
0 cat long	1.075770	-0.109050
dog short	1.643563	-1.469388
1 cat long	0.357021	-0.674600
dog short	-1.776904	-0.968914
2 cat long	-1.294524	0.413738
dog short	0.276662	-0.472035
3 cat long	-0.013960	-0.362543
dog short	-0.006154	-0.923061

The list of levels can contain either level names or level numbers (but not a mixture of the two).

```
# df.stack(level=['animal', 'hair_length'])
```

```
# from above is equivalent to:
```

```
In [27]: df.stack(level=[1, 2])
```

```
Out[27]:
```

exp	A	B
animal hair_length		
0 cat long	1.075770	-0.109050
dog short	1.643563	-1.469388
1 cat long	0.357021	-0.674600
dog short	-1.776904	-0.968914
2 cat long	-1.294524	0.413738
dog short	0.276662	-0.472035
3 cat long	-0.013960	-0.362543
dog short	-0.006154	-0.923061

Missing data

These functions are intelligent about handling missing data and do not expect each

subgroup within the hierarchical index to have the same set of labels. They also can handle the index being unsorted (but you can make it sorted by calling `sort_index`, of course). Here is a more complex example:

```
In [28]: columns = pd.MultiIndex.from_tuples([('A', 'cat'), ('B', 'dog'),
.....:                                     ('B', 'cat'), ('A', 'dog')],
.....:                                     names=['exp', 'animal'])
.....:
```

```
In [29]: index = pd.MultiIndex.from_product([('bar', 'baz', 'foo', 'qux'),
.....:                                     ('one', 'two')],
.....:                                     names=['first', 'second'])
.....:
```

```
In [30]: df = pd.DataFrame(np.random.randn(8, 4), index=index, columns=columns)
```

```
In [31]: df2 = df.iloc[[0, 1, 2, 4, 5, 7]]
```

```
In [32]: df2
```

```
Out[32]:
```

exp		A	B		A
animal		cat	dog	cat	dog
first	second				
bar	one	0.895717	0.805244	-1.206412	2.565646
	two	1.431256	1.340309	-1.170299	-0.226169
baz	one	0.410835	0.813850	0.132003	-0.827317
foo	one	-1.413681	1.607920	1.024180	0.569605
	two	0.875906	-2.211372	0.974466	-2.006747
qux	two	-1.226825	0.769804	-1.281247	-0.727707

As mentioned above, `stack` can be called with a `level` argument to select which level in the columns to stack:

```

In [33]: df2.stack('exp')
Out[33]:
animal      cat    dog
first second exp
bar one   A   0.895717 2.565646
        B  -1.206412 0.805244
      two   A   1.431256 -0.226169
        B  -1.170299 1.340309
baz one   A   0.410835 -0.827317
        B   0.132003 0.813850
foo one   A  -1.413681 0.569605
        B   1.024180 1.607920
      two   A   0.875906 -2.006747
        B   0.974466 -2.211372
qux two   A  -1.226825 -0.727707
        B  -1.281247 0.769804

```

```

In [34]: df2.stack('animal')
Out[34]:
exp      A      B
first second animal
bar one   cat   0.895717 -1.206412
        dog   2.565646 0.805244
      two   cat   1.431256 -1.170299
        dog  -0.226169 1.340309
baz one   cat   0.410835 0.132003
        dog  -0.827317 0.813850
foo one   cat  -1.413681 1.024180
        dog   0.569605 1.607920
      two   cat   0.875906 0.974466
        dog  -2.006747 -2.211372
qux two   cat  -1.226825 -1.281247
        dog  -0.727707 0.769804

```

Unstacking can result in missing values if subgroups do not have the same set of labels. By default, missing values will be replaced with the default fill value for that data type, `NaN` for float, `NaT` for datetimelike, etc. For integer types, by default data will be converted to float and missing values will be set to `NaN`.

```
In [35]: df3 = df.iloc[[0, 1, 4, 7], [1, 2]]
```

```
In [36]: df3
```

```
Out[36]:
```

```
exp      B
animal    dog    cat
first second
bar one    0.805244 -1.206412
    two    1.340309 -1.170299
foo one    1.607920 1.024180
qux two    0.769804 -1.281247
```

```
In [37]: df3.unstack()
```

```
Out[37]:
```

```
exp      B
animal    dog      cat
second    one    two    one    two
first
bar    0.805244 1.340309 -1.206412 -1.170299
foo    1.607920    NaN 1.024180    NaN
qux     NaN 0.769804    NaN -1.281247
```

New in version 0.18.0.

Alternatively, `unstack` takes an optional `fill_value` argument, for specifying the value of missing data.

```
In [38]: df3.unstack(fill_value=-1e9)
```

```
Out[38]:
```

```
exp      B
animal    dog      cat
second    one    two    one    two
first
bar    8.052440e-01 1.340309e+00 -1.206412e+00 -1.170299e+00
foo    1.607920e+00 -1.000000e+09 1.024180e+00 -1.000000e+09
qux    -1.000000e+09 7.698036e-01 -1.000000e+09 -1.281247e+00
```

With a MultiIndex

Unstacking when the columns are a `MultiIndex` is also careful about doing the right thing:

```
In [39]: df[:3].unstack(0)
Out[39]:
exp      A      B      A
animal   cat    dog    cat    dog
first   bar   baz   bar   baz   bar   baz   bar   baz
second
one    0.895717 0.410835 0.805244 0.81385 -1.206412 0.132003 2.565646 -0.827317
two    1.431256   NaN 1.340309   NaN -1.170299   NaN -0.226169   NaN
```

```
In [40]: df2.unstack(1)
Out[40]:
exp      A      B      A
animal   cat    dog    cat    dog
second   one   two   one   two   one   two   one   two
first
bar    0.895717 1.431256 0.805244 1.340309 -1.206412 -1.170299 2.565646 -0.226169
baz    0.410835   NaN 0.813850   NaN 0.132003   NaN -0.827317   NaN
foo   -1.413681 0.875906 1.607920 -2.211372 1.024180 0.974466 0.569605 -2.006747
qux     NaN -1.226825   NaN 0.769804   NaN -1.281247   NaN -0.727707
```

Reshaping by Melt

Melt

df3					df3.melt(id_vars=['first', 'last'])				
	first	last	height	weight		first	last	variable	value
0	John	Doe	5.5	130	0	John	Doe	height	5.5
1	Mary	Bo	6.0	150	1	Mary	Bo	height	6.0
					2	John	Doe	weight	130
					3	Mary	Bo	weight	150

The top-level `melt()` function and the corresponding `DataFrame.melt()` are useful to massage a `DataFrame` into a format where one or more columns are *identifier variables*, while all other columns, considered *measured variables*, are “unpivoted” to the row axis, leaving just two non-identifier columns, “variable” and “value”. The names of those columns can be customized by supplying the `var_name` and `value_name` parameters.

For instance,

```
In [41]: cheese = pd.DataFrame({'first': ['John', 'Mary'],
.....:                        'last': ['Doe', 'Bo'],
.....:                        'height': [5.5, 6.0],
.....:                        'weight': [130, 150]})
.....:
```

```
In [42]: cheese
```

```
Out[42]:
```

	first	last	height	weight
0	John	Doe	5.5	130
1	Mary	Bo	6.0	150

```
In [43]: cheese.melt(id_vars=['first', 'last'])
```

```
Out[43]:
```

	first	last	variable	value
0	John	Doe	height	5.5
1	Mary	Bo	height	6.0
2	John	Doe	weight	130.0
3	Mary	Bo	weight	150.0

```
In [44]: cheese.melt(id_vars=['first', 'last'], var_name='quantity')
```

```
Out[44]:
```

	first	last	quantity	value
0	John	Doe	height	5.5
1	Mary	Bo	height	6.0
2	John	Doe	weight	130.0
3	Mary	Bo	weight	150.0

Another way to transform is to use the [wide_to_long\(\)](#) panel data convenience function. It is less flexible than [melt\(\)](#), but more user-friendly.

```
In [45]: dft = pd.DataFrame({"A1970": {0: "a", 1: "b", 2: "c"},
.....:                     "A1980": {0: "d", 1: "e", 2: "f"},
.....:                     "B1970": {0: 2.5, 1: 1.2, 2: .7},
.....:                     "B1980": {0: 3.2, 1: 1.3, 2: .1},
.....:                     "X": dict(zip(range(3), np.random.randn(3)))
.....:                     })
.....:
```

```
In [46]: dft["id"] = dft.index
```

```
In [47]: dft
```

```
Out[47]:
```

```
   A1970 A1980 B1970 B1980      X id
0     a     d   2.5   3.2 -0.121306 0
1     b     e   1.2   1.3 -0.097883 1
2     c     f   0.7   0.1  0.695775 2
```

```
In [48]: pd.wide_to_long(dft, ["A", "B"], i="id", j="year")
```

```
Out[48]:
```

```
      X A  B
id year
0 1970 -0.121306 a 2.5
1 1970 -0.097883 b 1.2
2 1970  0.695775 c 0.7
0 1980 -0.121306 d 3.2
1 1980 -0.097883 e 1.3
2 1980  0.695775 f 0.1
```

Combining with stats and GroupBy

It should be no shock that combining `pivot` / `stack` / `unstack` with `GroupBy` and the basic `Series` and `DataFrame` statistical functions can produce some very expressive and fast data manipulations.

```
In [49]: df
Out[49]:
exp      A      B      A
animal   cat   dog   cat   dog
first second
bar one   0.895717 0.805244 -1.206412 2.565646
    two   1.431256 1.340309 -1.170299 -0.226169
baz one   0.410835 0.813850 0.132003 -0.827317
    two  -0.076467 -1.187678 1.130127 -1.436737
foo one  -1.413681 1.607920 1.024180 0.569605
    two   0.875906 -2.211372 0.974466 -2.006747
qux one  -0.410001 -0.078638 0.545952 -1.219217
    two  -1.226825 0.769804 -1.281247 -0.727707
```

```
In [50]: df.stack().mean(1).unstack()
```

```
Out[50]:
animal      cat      dog
first second
bar one  -0.155347 1.685445
    two   0.130479 0.557070
baz one   0.271419 -0.006733
    two   0.526830 -1.312207
foo one  -0.194750 1.088763
    two   0.925186 -2.109060
qux one   0.067976 -0.648927
    two  -1.254036 0.021048
```

```
# same result, another way
```

```
In [51]: df.groupby(level=1, axis=1).mean()
```

```
Out[51]:
animal      cat      dog
first second
bar one  -0.155347 1.685445
    two   0.130479 0.557070
baz one   0.271419 -0.006733
    two   0.526830 -1.312207
foo one  -0.194750 1.088763
    two   0.925186 -2.109060
qux one   0.067976 -0.648927
    two  -1.254036 0.021048
```

```
In [52]: df.stack().groupby(level=1).mean()
```

```
Out[52]:
exp      A      B
second
one   0.071448 0.455513
two  -0.424186 -0.204486
```

```
In [53]: df.mean().unstack(0)
```

```
Out[53]:
exp      A      B
animal
cat   0.060843 0.018596
dog  -0.413580 0.232430
```


Pivot tables

While `pivot()` provides general purpose pivoting with various data types (strings, numerics, etc.), pandas also provides `pivot_table()` for pivoting with aggregation of numeric data.

The function `pivot_table()` can be used to create spreadsheet-style pivot tables. See the [cookbook](#) for some advanced strategies.

It takes a number of arguments:

- `data` : a DataFrame object.
- `values` : a column or a list of columns to aggregate.
- `index` : a column, Grouper, array which has the same length as data, or list of them. Keys to group by on the pivot table index. If an array is passed, it is being used as the same manner as column values.
- `columns` : a column, Grouper, array which has the same length as data, or list of them. Keys to group by on the pivot table column. If an array is passed, it is being used as the same manner as column values.
- `aggfunc` : function to use for aggregation, defaulting to `numpy.mean`.

Consider a data set like this:

In [54]: `import datetime`

```
In [55]: df = pd.DataFrame({'A': ['one', 'one', 'two', 'three'] * 6,  
.....:                    'B': ['A', 'B', 'C'] * 8,  
.....:                    'C': ['foo', 'foo', 'foo', 'bar', 'bar', 'bar'] * 4,  
.....:                    'D': np.random.randn(24),  
.....:                    'E': np.random.randn(24),  
.....:                    'F': [datetime.datetime(2013, i, 1) for i in range(1, 13)]  
.....:                    + [datetime.datetime(2013, i, 15) for i in range(1, 13)]})  
.....:
```

In [56]: `df`

Out[56]:

	A	B	C	D	E	F
0	one	A	foo	0.341734	-0.317441	2013-01-01
1	one	B	foo	0.959726	-1.236269	2013-02-01
2	two	C	foo	-1.110336	0.896171	2013-03-01
3	three	A	bar	-0.619976	-0.487602	2013-04-01
4	one	B	bar	0.149748	-0.082240	2013-05-01
..
19	three	B	foo	0.690579	-2.213588	2013-08-15
20	one	C	foo	0.995761	1.063327	2013-09-15
21	one	A	bar	2.396780	1.266143	2013-10-15
22	two	B	bar	0.014871	0.299368	2013-11-15
23	three	C	bar	3.357427	-0.863838	2013-12-15

[24 rows x 6 columns]

We can produce pivot tables from this data very easily:

```
In [57]: pd.pivot_table(df, values='D', index=['A', 'B'], columns=['C'])
```

```
Out[57]:
```

```
C      bar      foo
A  B
one A  1.120915 -0.514058
    B -0.338421  0.002759
    C -0.538846  0.699535
three A -1.181568   NaN
      B   NaN  0.433512
      C  0.588783   NaN
two  A   NaN  1.000985
      B  0.158248   NaN
      C   NaN  0.176180
```

```
In [58]: pd.pivot_table(df, values='D', index=['B'], columns=['A', 'C'], aggfunc=np.sum)
```

```
Out[58]:
```

```
A      one      three      two
C      bar      foo      bar      foo      bar      foo
B
A  2.241830 -1.028115 -2.363137   NaN   NaN  2.001971
B -0.676843  0.005518   NaN  0.867024  0.316495   NaN
C -1.077692  1.399070  1.177566   NaN   NaN  0.352360
```

```
In [59]: pd.pivot_table(df, values=['D', 'E'], index=['B'], columns=['A', 'C'],
```

```
.....:      aggfunc=np.sum)
```

```
.....:
```

```
Out[59]:
```

```
      D      E
A      one      three      two      one      three      two
C      bar      foo      bar      foo      bar      foo      bar      foo
B
A  2.241830 -1.028115 -2.363137   NaN   NaN  2.001971  2.786113 -0.043211  1.922577
NaN   NaN  0.128491
B -0.676843  0.005518   NaN  0.867024  0.316495   NaN  1.368280 -1.103384   NaN -
2.128743 -0.194294   NaN
C -1.077692  1.399070  1.177566   NaN   NaN  0.352360 -1.976883  1.495717 -0.263660
NaN   NaN  0.872482
```

The result object is a **DataFrame** having potentially hierarchical indexes on the rows and columns. If the **values** column name is not given, the pivot table will include all of the data that can be aggregated in an additional level of hierarchy in the columns:

```
In [60]: pd.pivot_table(df, index=['A', 'B'], columns=['C'])
```

```
Out[60]:
```

		D		E	
C		bar	foo	bar	foo
A	B				
	one	A	1.120915	-0.514058	1.393057
		B	-0.338421	0.002759	0.684140
three		C	-0.538846	0.699535	-0.988442
	A	-1.181568	NaN	0.961289	NaN
	B	NaN	0.433512	NaN	-1.064372
two		C	0.588783	NaN	-0.131830
	A	NaN	1.000985	NaN	0.064245
	B	0.158248	NaN	-0.097147	NaN
	C	NaN	0.176180	NaN	0.436241

Also, you can use `Grouper` for `index` and `columns` keywords. For detail of `Grouper`, see [Grouping with a Grouper specification](#).

```
In [61]: pd.pivot_table(df, values='D', index=pd.Grouper(freq='M', key='F'),
```

```
.....:         columns='C')
```

```
.....:
```

```
Out[61]:
```

C	bar	foo
F		
2013-01-31	NaN	-0.514058
2013-02-28	NaN	0.002759
2013-03-31	NaN	0.176180
2013-04-30	-1.181568	NaN
2013-05-31	-0.338421	NaN
2013-06-30	-0.538846	NaN
2013-07-31	NaN	1.000985
2013-08-31	NaN	0.433512
2013-09-30	NaN	0.699535
2013-10-31	1.120915	NaN
2013-11-30	0.158248	NaN
2013-12-31	0.588783	NaN

You can render a nice output of the table omitting the missing values by calling `to_string` if you wish:

```
In [62]: table = pd.pivot_table(df, index=['A', 'B'], columns=['C'])
```

```
In [63]: print(table.to_string(na_rep=''))
```

		D		E	
C		bar	foo	bar	foo
A	B				
one	A	1.120915	-0.514058	1.393057	-0.021605
	B	-0.338421	0.002759	0.684140	-0.551692
	C	-0.538846	0.699535	-0.988442	0.747859
three	A	-1.181568		0.961289	
	B		0.433512		-1.064372
	C	0.588783		-0.131830	
two	A		1.000985		0.064245
	B	0.158248		-0.097147	
	C		0.176180		0.436241

Note that `pivot_table` is also available as an instance method on `DataFrame`, i.e. `DataFrame.pivot_table()`.

Adding margins

If you pass `margins=True` to `pivot_table`, special `All` columns and rows will be added with partial group aggregates across the categories on the rows and columns:

```
In [64]: df.pivot_table(index=['A', 'B'], columns='C', margins=True, aggfunc=np.std)
```

```
Out[64]:
```

		D		E			
C		bar	foo	All	bar	foo	All
A	B						
one	A	1.804346	1.210272	1.569879	0.179483	0.418374	0.858005
	B	0.690376	1.353355	0.898998	1.083825	0.968138	1.101401
	C	0.273641	0.418926	0.771139	1.689271	0.446140	1.422136
three	A	0.794212	NaN	0.794212	2.049040	NaN	2.049040
	B	NaN	0.363548	0.363548	NaN	1.625237	1.625237
	C	3.915454	NaN	3.915454	1.035215	NaN	1.035215
two	A	NaN	0.442998	0.442998	NaN	0.447104	0.447104
	B	0.202765	NaN	0.202765	0.560757	NaN	0.560757
	C	NaN	1.819408	1.819408	NaN	0.650439	0.650439
All		1.556686	0.952552	1.246608	1.250924	0.899904	1.059389

Cross tabulations

Use `crosstab()` to compute a cross-tabulation of two (or more) factors. By default `crosstab` computes a frequency table of the factors unless an array of values and an aggregation function are passed.

It takes a number of arguments

- `index` : array-like, values to group by in the rows.
- `columns` : array-like, values to group by in the columns.
- `values` : array-like, optional, array of values to aggregate according to the factors.

- `aggfunc` : function, optional, If no values array is passed, computes a frequency table.
- `rownames` : sequence, default `None` , must match number of row arrays passed.
- `colnames` : sequence, default `None` , if passed, must match number of column arrays passed.
- `margins` : boolean, default `False` , Add row/column margins (subtotals)
- `normalize` : boolean, {'all', 'index', 'columns'}, or {0,1}, default `False` . Normalize by dividing all values by the sum of values.

Any `Series` passed will have their name attributes used unless row or column names for the cross-tabulation are specified

For example:

```
In [65]: foo, bar, dull, shiny, one, two = 'foo', 'bar', 'dull', 'shiny', 'one', 'two'
```

```
In [66]: a = np.array([foo, foo, bar, bar, foo, foo], dtype=object)
```

```
In [67]: b = np.array([one, one, two, one, two, one], dtype=object)
```

```
In [68]: c = np.array([dull, dull, shiny, dull, dull, shiny], dtype=object)
```

```
In [69]: pd.crosstab(a, [b, c], rownames=['a'], colnames=['b', 'c'])
```

```
Out[69]:
```

```
b  one    two
c  dull shiny dull shiny
a
bar  1    0    0    1
foo  2    1    1    0
```

If `crosstab` receives only two Series, it will provide a frequency table.

```
In [70]: df = pd.DataFrame({'A': [1, 2, 2, 2, 2], 'B': [3, 3, 4, 4, 4],
.....:                    'C': [1, 1, np.nan, 1, 1]})
.....:
```

```
In [71]: df
```

```
Out[71]:
```

```
  A  B  C
0  1  3  1.0
1  2  3  1.0
2  2  4  NaN
3  2  4  1.0
4  2  4  1.0
```

```
In [72]: pd.crosstab(df.A, df.B)
```

```
Out[72]:
```

```
B  3  4
A
1  1  0
2  1  3
```

Any input passed containing **Categorical** data will have **all** of its categories included in the cross-tabulation, even if the actual data does not contain any instances of a particular category.

```
In [73]: foo = pd.Categorical(['a', 'b'], categories=['a', 'b', 'c'])
```

```
In [74]: bar = pd.Categorical(['d', 'e'], categories=['d', 'e', 'f'])
```

```
In [75]: pd.crosstab(foo, bar)
```

```
Out[75]:
```

```
col_0 d e
row_0
a     1  0
b     0  1
```

Normalization

New in version 0.18.1.

Frequency tables can also be normalized to show percentages rather than counts using the **normalize** argument:

```
In [76]: pd.crosstab(df.A, df.B, normalize=True)
```

```
Out[76]:
```

```
B   3   4
A
1  0.2  0.0
2  0.2  0.6
```

normalize can also normalize values within each row or within each column:

```
In [77]: pd.crosstab(df.A, df.B, normalize='columns')
```

```
Out[77]:
```

```
B   3   4
A
1  0.5  0.0
2  0.5  1.0
```

crosstab can also be passed a third **Series** and an aggregation function (**aggfunc**) that will be applied to the values of the third **Series** within each group defined by the first two **Series** :

```
In [78]: pd.crosstab(df.A, df.B, values=df.C, aggfunc=np.sum)
```

```
Out[78]:
```

```
B   3   4
A
1  1.0 NaN
2  1.0  2.0
```

Adding margins

Finally, one can also add margins or normalize this output.

```
In [79]: pd.crosstab(df.A, df.B, values=df.C, aggfunc=np.sum, normalize=True,
....:               margins=True)
....:
Out[79]:
B    3    4  All
A
1   0.25  0.0  0.25
2   0.25  0.5  0.75
All  0.50  0.5  1.00
```

Tiling

The `cut()` function computes groupings for the values of the input array and is often used to transform continuous variables to discrete or categorical variables:

```
In [80]: ages = np.array([10, 15, 13, 12, 23, 25, 28, 59, 60])
```

```
In [81]: pd.cut(ages, bins=3)
```

```
Out[81]:
```

```
[(9.95, 26.667], (9.95, 26.667], (9.95, 26.667], (9.95, 26.667], (9.95, 26.667], (9.95, 26.667], (26.667,
43.333], (43.333, 60.0], (43.333, 60.0]]
```

```
Categories (3, interval[float64]): [(9.95, 26.667] < (26.667, 43.333] < (43.333, 60.0]]
```

If the `bins` keyword is an integer, then equal-width bins are formed. Alternatively we can specify custom bin-edges:

```
In [82]: c = pd.cut(ages, bins=[0, 18, 35, 70])
```

```
In [83]: c
```

```
Out[83]:
```

```
[(0, 18], (0, 18], (0, 18], (0, 18], (18, 35], (18, 35], (18, 35], (35, 70], (35, 70]]
```

```
Categories (3, interval[int64]): [(0, 18] < (18, 35] < (35, 70]]
```

New in version 0.20.0.

If the `bins` keyword is an `IntervalIndex`, then these will be used to bin the passed data.:

```
pd.cut([25, 20, 50], bins=c.categories)
```

Computing indicator / dummy variables

To convert a categorical variable into a “dummy” or “indicator” `DataFrame`, for example a column in a `DataFrame` (a `Series`) which has `k` distinct values, can derive a `DataFrame` containing `k` columns of 1s and 0s using `get_dummies()`:

```
In [84]: df = pd.DataFrame({'key': list('bbacab'), 'data1': range(6)})
```

```
In [85]: pd.get_dummies(df['key'])
```

```
Out[85]:
```

	a	b	c
0	0	1	0
1	0	1	0
2	1	0	0
3	0	0	1
4	1	0	0
5	0	1	0

Sometimes it's useful to prefix the column names, for example when merging the result with the original `DataFrame` :

```
In [86]: dummies = pd.get_dummies(df['key'], prefix='key')
```

```
In [87]: dummies
```

```
Out[87]:
```

	key_a	key_b	key_c
0	0	1	0
1	0	1	0
2	1	0	0
3	0	0	1
4	1	0	0
5	0	1	0

```
In [88]: df[['data1']].join(dummies)
```

```
Out[88]:
```

	data1	key_a	key_b	key_c
0	0	0	1	0
1	1	0	1	0
2	2	1	0	0
3	3	0	0	1
4	4	1	0	0
5	5	0	1	0

This function is often used along with discretization functions like `cut` :


```
In [89]: values = np.random.randn(10)
```

```
In [90]: values
```

```
Out[90]:
```

```
array([ 0.4082, -1.0481, -0.0257, -0.9884,  0.0941,  1.2627,  1.29 ,
        0.0824, -0.0558,  0.5366])
```

```
In [91]: bins = [0, 0.2, 0.4, 0.6, 0.8, 1]
```

```
In [92]: pd.get_dummies(pd.cut(values, bins))
```

```
Out[92]:
```

	(0.0, 0.2]	(0.2, 0.4]	(0.4, 0.6]	(0.6, 0.8]	(0.8, 1.0]
0	0	0	1	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	1	0	0	0	0
5	0	0	0	0	0
6	0	0	0	0	0
7	1	0	0	0	0
8	0	0	0	0	0
9	0	0	1	0	0

See also [Series.str.get_dummies](#).

[get_dummies\(\)](#) also accepts a `DataFrame`. By default all categorical variables (categorical in the statistical sense, those with *object* or *categorical* dtype) are encoded as dummy variables.

```
In [93]: df = pd.DataFrame({'A': ['a', 'b', 'a'], 'B': ['c', 'c', 'b'],
.....:                    'C': [1, 2, 3]})
.....:
```

```
In [94]: pd.get_dummies(df)
```

```
Out[94]:
```

	C	A_a	A_b	B_b	B_c
0	1	1	0	0	1
1	2	0	1	0	1
2	3	1	0	1	0

All non-object columns are included untouched in the output. You can control the columns that are encoded with the `columns` keyword.

```
In [95]: pd.get_dummies(df, columns=['A'])
```

```
Out[95]:
```

	B	C	A_a	A_b
0	c	1	1	0
1	c	2	0	1
2	b	3	1	0

Notice that the `B` column is still included in the output, it just hasn't been encoded. You can drop `B` before calling `get_dummies` if you don't want to include it in the output.

As with the `Series` version, you can pass values for the `prefix` and `prefix_sep`. By default the column name is used as the prefix, and `'_'` as the prefix separator. You can specify `prefix` and `prefix_sep` in 3 ways:

- string: Use the same value for `prefix` or `prefix_sep` for each column to be encoded.
- list: Must be the same length as the number of columns being encoded.
- dict: Mapping column name to prefix.

```
In [96]: simple = pd.get_dummies(df, prefix='new_prefix')
```

```
In [97]: simple
```

```
Out[97]:
```

	C	new_prefix_a	new_prefix_b	new_prefix_b	new_prefix_c
0	1	1	0	0	1
1	2	0	1	0	1
2	3	1	0	1	0

```
In [98]: from_list = pd.get_dummies(df, prefix=['from_A', 'from_B'])
```

```
In [99]: from_list
```

```
Out[99]:
```

	C	from_A_a	from_A_b	from_B_b	from_B_c
0	1	1	0	0	1
1	2	0	1	0	1
2	3	1	0	1	0

```
In [100]: from_dict = pd.get_dummies(df, prefix={'B': 'from_B', 'A': 'from_A'})
```

```
In [101]: from_dict
```

```
Out[101]:
```

	C	from_A_a	from_A_b	from_B_b	from_B_c
0	1	1	0	0	1
1	2	0	1	0	1
2	3	1	0	1	0

New in version 0.18.0.

Sometimes it will be useful to only keep $k-1$ levels of a categorical variable to avoid collinearity when feeding the result to statistical models. You can switch to this mode by turn on `drop_first`.

```
In [102]: s = pd.Series(list('abcaa'))
```

```
In [103]: pd.get_dummies(s)
```

```
Out[103]:
```

```
   a b c
0  1 0 0
1  0 1 0
2  0 0 1
3  1 0 0
4  1 0 0
```

```
In [104]: pd.get_dummies(s, drop_first=True)
```

```
Out[104]:
```

```
   b c
0  0 0
1  1 0
2  0 1
3  0 0
4  0 0
```

When a column contains only one level, it will be omitted in the result.

```
In [105]: df = pd.DataFrame({'A': list('aaaaa'), 'B': list('ababc')})
```

```
In [106]: pd.get_dummies(df)
```

```
Out[106]:
```

```
   A_a B_a B_b B_c
0    1    1    0    0
1    1    0    1    0
2    1    1    0    0
3    1    0    1    0
4    1    0    0    1
```

```
In [107]: pd.get_dummies(df, drop_first=True)
```

```
Out[107]:
```

```
   B_b B_c
0    0    0
1    1    0
2    0    0
3    1    0
4    0    1
```

By default new columns will have `np.uint8` dtype. To choose another dtype, use the `dtype` argument:

```
In [108]: df = pd.DataFrame({'A': list('abc'), 'B': [1.1, 2.2, 3.3]})
```

```
In [109]: pd.get_dummies(df, dtype=bool).dtypes
```

```
Out[109]:
```

```
B      float64
A_a      bool
A_b      bool
A_c      bool
dtype: object
```

New in version 0.23.0.

Factorizing values

To encode 1-d values as an enumerated type use `factorize()`:

```
In [110]: x = pd.Series(['A', 'A', np.nan, 'B', 3.14, np.inf])
```

```
In [111]: x
```

```
Out[111]:
```

```
0    A
1    A
2   NaN
3    B
4   3.14
5   inf
```

```
dtype: object
```

```
In [112]: labels, uniques = pd.factorize(x)
```

```
In [113]: labels
```

```
Out[113]: array([ 0,  0, -1,  1,  2,  3])
```

```
In [114]: uniques
```

```
Out[114]: Index(['A', 'B', 3.14, inf], dtype='object')
```

Note that `factorize` is similar to `numpy.unique` , but differs in its handling of NaN:

Note

The following `numpy.unique` will fail under Python 3 with a `TypeError` because of an ordering bug. See also [here](#).

```
In [1]: x = pd.Series(['A', 'A', np.nan, 'B', 3.14, np.inf])
```

```
In [2]: pd.factorize(x, sort=True)
```

```
Out[2]:
```

```
(array([ 2,  2, -1,  3,  0,  1]),
 Index([3.14, inf, 'A', 'B'], dtype='object'))
```

```
In [3]: np.unique(x, return_inverse=True)[::-1]
```

```
Out[3]: (array([3, 3, 0, 4, 1, 2]), array([nan, 3.14, inf, 'A', 'B'], dtype=object))
```

Note

If you just want to handle one column as a categorical variable (like R's factor), you can use `df["cat_col"] = pd.Categorical(df["col"])` or `df["cat_col"] = df["col"].astype("category")` . For full docs on [Categorical](#), see the [Categorical introduction](#) and the [API documentation](#).

Examples

In this section, we will review frequently asked questions and examples. The column names and relevant column values are named to correspond with how this DataFrame will be pivoted in the answers below.

```
In [115]: np.random.seed([3, 1415])
```

```
In [116]: n = 20
```

```
In [117]: cols = np.array(['key', 'row', 'item', 'col'])
```

```
In [118]: df = cols + pd.DataFrame((np.random.randint(5, size=(n, 4))
.....:                             // [2, 1, 2, 1]).astype(str))
.....:
```

```
In [119]: df.columns = cols
```

```
In [120]: df = df.join(pd.DataFrame(np.random.rand(n, 2).round(2)).add_prefix('val'))
```

```
In [121]: df
```

```
Out[121]:
```

```
   key row item col val0 val1
0 key0 row3 item1 col3 0.81 0.04
1 key1 row2 item1 col2 0.44 0.07
2 key1 row0 item1 col0 0.77 0.01
3 key0 row4 item0 col2 0.15 0.59
4 key1 row0 item2 col1 0.81 0.64
..  ...
15 key0 row3 item1 col1 0.31 0.23
16 key0 row0 item2 col3 0.86 0.01
17 key0 row4 item0 col3 0.64 0.21
18 key2 row2 item2 col0 0.13 0.45
19 key0 row2 item0 col4 0.37 0.70
```

```
[20 rows x 6 columns]
```

Pivoting with single aggregations

Suppose we wanted to pivot `df` such that the `col` values are columns, `row` values are the index, and the mean of `val0` are the values? In particular, the resulting DataFrame should look like:

Note

```
col col0 col1 col2 col3 col4 row row0 0.77 0.605 NaN 0.860 0.65 row2 0.13 NaN 0.395
0.500 0.25 row3 NaN 0.310 NaN 0.545 NaN row4 NaN 0.100 0.395 0.760 0.24
```

This solution uses `pivot_table()`. Also note that `aggfunc='mean'` is the default. It is included here to be explicit.

```
In [122]: df.pivot_table(
.....:     values='val0', index='row', columns='col', aggfunc='mean')
.....:
Out[122]:
col col0 col1 col2 col3 col4
row
row0 0.77 0.605 NaN 0.860 0.65
row2 0.13 NaN 0.395 0.500 0.25
row3 NaN 0.310 NaN 0.545 NaN
row4 NaN 0.100 0.395 0.760 0.24
```

Note that we can also replace the missing values by using the `fill_value` parameter.

```
In [123]: df.pivot_table(
.....:     values='val0', index='row', columns='col', aggfunc='mean', fill_value=0)
.....:
Out[123]:
col col0 col1 col2 col3 col4
row
row0 0.77 0.605 0.000 0.860 0.65
row2 0.13 0.000 0.395 0.500 0.25
row3 0.00 0.310 0.000 0.545 0.00
row4 0.00 0.100 0.395 0.760 0.24
```

Also note that we can pass in other aggregation functions as well. For example, we can also pass in `sum`.

```
In [124]: df.pivot_table(
.....:     values='val0', index='row', columns='col', aggfunc='sum', fill_value=0)
.....:
Out[124]:
col col0 col1 col2 col3 col4
row
row0 0.77 1.21 0.00 0.86 0.65
row2 0.13 0.00 0.79 0.50 0.50
row3 0.00 0.31 0.00 1.09 0.00
row4 0.00 0.10 0.79 1.52 0.24
```

Another aggregation we can do is calculate the frequency in which the columns and rows occur together a.k.a. “cross tabulation”. To do this, we can pass `size` to the `aggfunc` parameter.

```
In [125]: df.pivot_table(index='row', columns='col', fill_value=0, aggfunc='size')
Out[125]:
col col0 col1 col2 col3 col4
row
row0 1 2 0 1 1
row2 1 0 2 1 2
row3 0 1 0 2 0
row4 0 1 2 2 1
```

Pivoting with multiple aggregations

We can also perform multiple aggregations. For example, to perform both a `sum` and `mean`, we can pass in a list to the `aggfunc` argument.

```
In [126]: df.pivot_table(
.....:     values='val0', index='row', columns='col', aggfunc=['mean', 'sum'])
.....:
Out[126]:
      mean      sum
col col0 col1 col2 col3 col4 col0 col1 col2 col3 col4
row
row0 0.77 0.605 NaN 0.860 0.65 0.77 1.21 NaN 0.86 0.65
row2 0.13 NaN 0.395 0.500 0.25 0.13 NaN 0.79 0.50 0.50
row3 NaN 0.310 NaN 0.545 NaN NaN 0.31 NaN 1.09 NaN
row4 NaN 0.100 0.395 0.760 0.24 NaN 0.10 0.79 1.52 0.24
```

Note to aggregate over multiple value columns, we can pass in a list to the `values` parameter.

```
In [127]: df.pivot_table(
.....:     values=['val0', 'val1'], index='row', columns='col', aggfunc=['mean'])
.....:
Out[127]:
      mean      val1
col col0 col1 col2 col3 col4 col0 col1 col2 col3 col4
row
row0 0.77 0.605 NaN 0.860 0.65 0.01 0.745 NaN 0.010 0.02
row2 0.13 NaN 0.395 0.500 0.25 0.45 NaN 0.34 0.440 0.79
row3 NaN 0.310 NaN 0.545 NaN NaN 0.230 NaN 0.075 NaN
row4 NaN 0.100 0.395 0.760 0.24 NaN 0.070 0.42 0.300 0.46
```

Note to subdivide over multiple columns we can pass in a list to the `columns` parameter.

```
In [128]: df.pivot_table(
.....:     values=['val0'], index='row', columns=['item', 'col'], aggfunc=['mean'])
.....:
Out[128]:
      mean
      val0
item item0      item1      item2
col col2 col3 col4 col0 col1 col2 col3 col4 col0 col1 col3 col4
row
row0 NaN NaN NaN 0.77 NaN NaN NaN NaN NaN NaN 0.605 0.86 0.65
row2 0.35 NaN 0.37 NaN NaN 0.44 NaN NaN 0.13 NaN 0.50 0.13
row3 NaN NaN NaN NaN 0.31 NaN 0.81 NaN NaN NaN 0.28 NaN
row4 0.15 0.64 NaN NaN 0.10 0.64 0.88 0.24 NaN NaN NaN NaN
```

Exploding a list-like column

New in version 0.25.0.

Sometimes the values in a column are list-like.

```
In [129]: keys = ['panda1', 'panda2', 'panda3']
```

```
In [130]: values = [['eats', 'shoots'], ['shoots', 'leaves'], ['eats', 'leaves']]
```

```
In [131]: df = pd.DataFrame({'keys': keys, 'values': values})
```

```
In [132]: df
```

```
Out[132]:
```

	keys	values
0	panda1	[eats, shoots]
1	panda2	[shoots, leaves]
2	panda3	[eats, leaves]

We can 'explode' the `values` column, transforming each list-like to a separate row, by using `explode()`. This will replicate the index values from the original row:

```
In [133]: df['values'].explode()
```

```
Out[133]:
```

0	eats
0	shoots
1	shoots
1	leaves
2	eats
2	leaves

```
Name: values, dtype: object
```

You can also explode the column in the `DataFrame` .

```
In [134]: df.explode('values')
```

```
Out[134]:
```

	keys	values
0	panda1	eats
0	panda1	shoots
1	panda2	shoots
1	panda2	leaves
2	panda3	eats
2	panda3	leaves

`Series.explode()` will replace empty lists with `np.nan` and preserve scalar entries. The dtype of the resulting `Series` is always `object` .


```
In [135]: s = pd.Series([[1, 2, 3], 'foo', [], ['a', 'b']])
```

```
In [136]: s
```

```
Out[136]:
```

```
0    [1, 2, 3]
1         foo
2          []
3     [a, b]
dtype: object
```

```
In [137]: s.explode()
```

```
Out[137]:
```

```
0    1
0    2
0    3
1   foo
2   NaN
3    a
3    b
dtype: object
```

Here is a typical usecase. You have comma separated strings in a column and want to expand this.

```
In [138]: df = pd.DataFrame([{'var1': 'a,b,c', 'var2': 1},
.....:                      {'var1': 'd,e,f', 'var2': 2}])
.....:
```

```
In [139]: df
```

```
Out[139]:
```

```
   var1  var2
0  a,b,c    1
1  d,e,f    2
```

Creating a long form DataFrame is now straightforward using explode and chained operations

```
In [140]: df.assign(var1=df.var1.str.split(',')).explode('var1')
```

```
Out[140]:
```

```
   var1  var2
0    a     1
0    b     1
0    c     1
1    d     2
1    e     2
1    f     2
```